Session 2: Iterate, Critique, Trust

Business Intelligence & AI

Kerry Back, Rice University

Homework Debrief

What worked? What didn’t?

Share Your Best and Worst

Your Best Output

- What task did you give Claude?
- How good was the first draft?
- How much editing did it need?
- How much time did it save?

Your Worst Output

- What went wrong?
- Was it factually incorrect?
- Was it generic or off-target?
- What would have fixed it?

Patterns from Your Experience

Where did AI save time? Where did it need heavy editing?

Today’s Agenda

Iterate Turn mediocre drafts into great ones

→

Trust Calibrate verification to stakes

→

Skills Build reusable prompts and workflows

→

Build Create your own recurring workflow

The Key Principle

Final Draft, Not First Draft

> Don’t review AI’s **first draft**. Instead:
>
> 1. Ask AI to produce the work
> 2. Ask it to **critique itself** and revise
> 3. Repeat until no further revisions are needed
> 4. **Now** review what it produced

You review a **final draft** instead of a first draft. This saves your time (though it burns tokens). The more consequential the work, the more revision cycles it deserves.

The Iteration Cycle

Draft Get the first version — fast but mediocre

→

Direct Give specific feedback, not "make it better"

→

Self-Critique Ask Claude to find its own weaknesses

→

Fix Apply the critique — dramatically better output

Demo Walkthrough

From mediocre draft to polished output

What Just Happened

The iteration sequence

Round 1: “This is too generic — make it specific to our Q4 results.” → Output improves from template to tailored.

Round 2: “The chart is misleading — use a zero baseline.” → Visual accuracy corrected.

Round 3: “Rewrite the executive summary in half the length.” → Concise, focused, executive-ready.

Round 4: “What’s weak about this analysis? What would a skeptical CFO push back on?” → Claude identifies 3-4 real gaps.

Round 5: “Fix those issues.” → Final version is dramatically stronger.

The Secret: Specificity

Vague (Weak)

- “Make it better”
- “Improve the analysis”
- “Make it more professional”
- “Fix the chart”
- “Add more detail”

Specific (Strong)

- “Lead with the biggest risk to Q4 revenue”
- “Add year-over-year comparisons for each metric”
- “Cut the background section and start with the recommendation”
- “Use a zero baseline and add data labels”
- “Include the three customers that drove the variance”

The Maker/Checker Framework

The Core Shift

The professional skill is shifting:

From producing analysis → **To** directing and verifying it

AI is the **maker**. You are the **checker**.

Calibrated Trust

Low Stakes

- Internal notes
- Brainstorming sessions
- Personal exploration
- First-draft ideas

Verification: Quick sanity check. Glance at axes, skim for obvious errors.

Medium Stakes

- Team presentations
- Internal reports
- Meeting summaries
- Analysis for your manager

Verification: Verify every number. Check that framing matches reality. Test key claims.

High Stakes

- Board presentations
- Regulatory filings
- Public-facing documents
- Financial disclosures

Verification: Independent verification of every claim. Cross-check sources. Second reviewer.

The Critical Question

What happens if this output is wrong?

If the answer is “nothing much” → quick check.

If the answer is “serious consequences” → full verification.

Demo: Skills and Structured Prompts

What Is a Skill?

> A **skill** is a reusable set of instructions that produces consistent output every time.
>
> Write it once. Reuse it every week. Get the same quality without re-explaining.

Think of it as a template for your AI — not a document template with blanks to fill in, but an instruction set that tells the AI exactly what to produce, in what format, with what tone and structure.

Example: Weekly Status Report Skill

Reusable prompt

```
You are a business analyst preparing a weekly status report.

Format: Start with a 2-sentence executive summary. Then three sections: 1. Key Accomplishments (bullet points, max 5) 2. Risks and Blockers (bullet points with severity: high/medium/low) 3. Next Week Priorities (numbered list, max 4)

Tone: Direct, factual, no filler. Write for a VP who reads 20 of these. Length: One page maximum. Always end with: “Decision needed:” followed by any open items requiring management input, or “None this week.”
```

The Critique Pattern: Three Independent Reviewers

Correctness & Completeness

Are there factual errors, logical gaps, or missing information? Are claims supported and assumptions justified?

Clarity & Persuasiveness

Is the structure logical? Is anything confusing, redundant, or buried? Could the main message be stated more directly?

Devil's Advocate

What are the strongest counterarguments? Where is the reasoning weakest? What would a skeptic challenge?

Three parallel reviewers examine the same output from different angles, then synthesize findings. The result is dramatically more reliable than any single-pass review.

Demo Walkthrough

Skills, structured prompts, and the /critique pattern

What You Just Saw

The workflow

1. **Skill created:** Reusable instructions for a recurring task
2. **Consistent output:** Same format, tone, and structure every time
3. **Sub-agent critique:** Three independent reviewers catch different issues
4. **Final version:** Incorporates all critique — stronger than any single pass

Hands-On: Build Your Own Workflow

Choose Your Workflow

Option 1: Report Template

- Pick a recurring report you write
- Define the sections, format, and tone
- Specify what data goes in each section
- Build a reusable prompt
- Test it with real or sample data

Option 2: Data Analysis

- Define a standard analysis pattern
- “Given a CSV, always produce these 5 charts and a summary”
- Specify chart types and formatting
- Include interpretation guidance
- Test with the sample dataset

Option 3: Document Review

- Build a review checklist
- “Read this document and flag issues in these categories”
- Specify severity levels
- Include tone and audience checks
- Test with a real document

Tips for Building Your Workflow

What makes a good skill

- **Be explicit about format:** sections, length, bullet vs. paragraph, headers
- **Specify tone:** formal vs. conversational, direct vs. diplomatic
- **Define the audience:** who reads this? What do they care about?
- **Include examples:** “A good bullet point looks like this: ...”
- **Add constraints:** “Never exceed one page” or “Always include a recommendation”

Test it 2-3 times with different inputs. A good skill works consistently, not just once.

Breakout: Your Productivity Portfolio

Groups of 3-4

Breakout Discussion Guide

Share and Compare

- What workflow did you build?
- What recurring task does it address?
- How much time does it save?
- What trust level does it require?

Identify Top Use Cases

- Which workflows save the most time?
- Which have acceptable risk levels?
- What patterns are common?
- **Each group:** pick your top 3 highest-value use cases

Group Reports

Each group: your top 3 highest-value AI use cases.

Bridge to Enterprise Data

The Shift Ahead

Sessions 1-2: You bring the data

- Copy-paste into Claude
- Upload documents
- Manual data entry
- Limited to what you can carry
- Great for personal productivity

Sessions 3+: AI goes to the data

- AI queries enterprise databases
- CRM, ERP, HR, finance systems
- Cross-system analysis in seconds
- The same skills — iteration, critique, verification
- Organizational-scale impact

Preview: XYZ Corp

A simulated enterprise environment

- **$500M B2B industrial supplies distributor** with three operating divisions
- **9 enterprise systems:** Salesforce, NetSuite, SAP, Workday, HubSpot, Oracle, and more
- **41 tables, 26,000+ rows, 3 years of data**
- Ask questions in plain English — the AI queries the databases for you
- Cross-system analysis: “Which customers buy from multiple divisions?”

Homework

Your Assignment This Week

Homework — Due Session 3

1. **Continue using Claude for real work.** Build at least one reusable workflow and test it 2-3 times. Refine until it produces consistent quality.
2. **Read the enterprise systems background document** (provided). It covers what Salesforce, SAP, Workday, NetSuite, and other enterprise systems actually do. Come to Session 3 with a basic vocabulary.
3. **Map your company's data landscape:**
    - What are your top 3-5 enterprise systems?
    - What cross-system questions can nobody answer easily today?
    - Bring your list to Session 3.

What Is Coming Next

Session 3

- AI meets enterprise data
- XYZ Corp hands-on
- Cross-system queries
- Your company's data map

Session 4

- Data to deliverables
- Full reporting pipelines
- The agent loop explained
- Build your own report

Sessions 5-6

- Deployment architecture
- Vendor evaluation
- Security and governance
- Risk frameworks

Sessions 7-8

- Predictive modeling (classification, regression, gradient boosting, model evaluation)
- Capstone strategy presentation
- Pitch to leadership